

- [14] Forough Arabshahi, Sameer Singh, and Animashree Anandkumar. 2018. Combining symbolic expressions and black-box function evaluations in neural programs. *arXiv preprint arXiv:1801.04342* (2018).
- [15] Shushan Arakelyan, Anna Hakhverdyan, Miltiadis Allamanis, Luis Garcia, Christophe Hauser, and Xiang Ren. 2022. NS3: Neuro-symbolic semantic code search. *Advances in Neural Information Processing Systems* 35 (2022), 10476–10491.
- [16] Luca Arrotta, Gabriele Civitarese, and Claudio Bettini. 2023. Neuro-Symbolic Approaches for Context-Aware Human Activity Recognition. *arXiv preprint arXiv:2306.05058* (2023).
- [17] Luca Arrotta, Gabriele Civitarese, and Claudio Bettini. 2024. Semantic Loss: A New Neuro-Symbolic Approach for Context-Aware Human Activity Recognition. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies* 7, 4 (2024), 1–29.
- [18] Masataro Asai and Christian Muise. 2020. Learning neural-symbolic descriptive planning models via cube-space priors: The voyage home (to STRIPS). *arXiv preprint arXiv:2004.12850* (2020).
- [19] Chace Ashcraft, Jennifer Sleeman, Caroline Tang, Jay Brett, and Anand Gnanadesikan. 2023. Neuro-Symbolic Bi-Directional Translation–Deep Learning Explainability for Climate Tipping Point Research. *arXiv preprint arXiv:2306.11161* (2023).
- [20] Yaniv Aspis, Krysia Broda, Jorge Lobo, and Alessandra Russo. 2022. Embed2sym-scalable neuro-symbolic reasoning via clustered embeddings. In *Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning*, Vol. 19. 421–431.
- [21] Bahador Bahmani, Hyoung Suk Suh, and WaiChing Sun. 2024. Discovering interpretable elastoplasticity models via the neural polynomial method enabled symbolic regressions. *Computer Methods in Applied Mechanics and Engineering* 422 (2024), 116827.
- [22] Jonathan Balloch, Zhiyu Lin, Robert Wright, Xiangyu Peng, Mustafa Hussain, Aaron Srinivas, Julia Kim, and Mark O Riedl. 2023. Neuro-Symbolic World Models for Adapting to Open World Novelty. *arXiv preprint arXiv:2301.06294* (2023).
- [23] Joanna Baran and Jan Kocoń. 2022. Linguistic knowledge application to neuro-symbolic transformers in sentiment analysis. In *2022 IEEE International Conference on Data Mining Workshops (ICDMW)*. IEEE, 395–402.
- [24] Pablo Barceló, Tamara Cucumides, Floris Geerts, Juan Reutter, and Miguel Romero. 2023. A neuro-symbolic framework for answering conjunctive queries. *arXiv preprint arXiv:2310.04598* (2023).
- [25] Kexin Gu Baugh, Nuri Cingillioglu, and Alessandra Russo. 2023. Neuro-symbolic Rule Learning in Real-world Classification Tasks. *arXiv preprint arXiv:2303.16674* (2023).
- [26] Antigoni Belekou, Charalabos Papageorgiou, Efstratios Karavasilis, Eleftheria Tsaltas, Nikolaos Kelekis, Christoph Klein, and Nikolaos Smyrnis. 2022. Paradoxical Reasoning: An fMRI Study. *Frontiers in Psychology* 13 (2022), 850491.
- [27] Paul Bello and Bertram F Malle. 2023. Computational Approaches to Morality. *The Cambridge Handbook of Computational Cognitive Sciences* 2 (2023), 1037–1063.
- [28] Tommaso Bendinelli, Luca Biggio, and Pierre-Alexandre Kamienny. 2023. Controllable neural symbolic regression. In *International Conference on Machine Learning*. PMLR, 2063–2077.
- [29] Adrien Bennetot, Jean-Luc Laurent, Raja Chatila, and Natalia Díaz-Rodríguez. 2019. Towards explainable neural-symbolic visual reasoning. *arXiv preprint arXiv:1909.09065* (2019).
- [30] Tarek R Besold, Artur d’Avila Garcez, Keith Stenning, Leendert van der Torre, and Michiel van Lambalgen. 2017. Reasoning in non-probabilistic uncertainty: Logic programming and neural-symbolic computing as examples. *Minds and Machines* 27 (2017), 37–77.
- [31] Luca Biggio, Tommaso Bendinelli, Alexander Neitz, Aurelien Lucchi, and Giambattista Parascandolo. 2021. Neural symbolic regression that scales. In *International Conference on Machine Learning*. Pmlr, 936–945.
- [32] Pierre Bonzon. 2017. Towards neuro-inspired symbolic models of cognition: linking neural dynamics to behaviors through asynchronous communications. *Cognitive neurodynamics* 11 (2017), 327–353.
- [33] Antoine Bosselut, Ronan Le Bras, and Yejin Choi. 2021. Dynamic neuro-symbolic knowledge graph construction for zero-shot commonsense question answering. In *Proceedings of the AAAI conference on Artificial Intelligence*, Vol. 35. 4923–4931.
- [34] David A Broniatowski et al. 2021. Psychological foundations of explainability and interpretability in artificial intelligence. *NIST, Tech. Rep* (2021).
- [35] Jonathon Cai, Richard Shin, and Dawn Song. 2017. Making neural programming architectures generalize via recursion. *arXiv preprint arXiv:1704.06611* (2017).
- [36] Roberta Calegari, Giovanni Ciatto, and Andrea Omicini. 2020. On the integration of symbolic and sub-symbolic techniques for XAI: A survey. *Intelligenza Artificiale* 14, 1 (2020), 7–32.
- [37] Tommaso Carraro. 2023. Overcoming Recommendation Limitations with Neuro-Symbolic Integration. In *Proceedings of the 17th ACM Conference on Recommender Systems*. 1325–1331.
- [38] Lupita Estefania Gazzo Castañeda, Benjamin Sklarek, Dennis E Dal Mas, and Markus Knauff. 2023. Probabilistic and deductive reasoning in the human brain. *NeuroImage* 275 (2023), 120180.
- [39] Subhajt Chaudhury, Prithviraj Sen, Masaki Ono, Daiki Kimura, Michiaki Tatsubori, and Asim Munawar. 2021. Neuro-symbolic approaches for text-based policy learning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 3073–3078.
- [40] Subhajt Chaudhury, Sarathkrishna Swaminathan, Daiki Kimura, Prithviraj Sen, Keerthiram Murugesan, Rosario Uceda-Sosa, Michiaki Tatsubori, Achille Fokoue, Pavan Kapanipathi, Asim Munawar, et al. 2023. Learning symbolic rules over abstract meaning representations for textual reinforcement learning. *arXiv preprint arXiv:2307.02689* (2023).
- [41] Subhajt Chaudhury, Sarathkrishna Swaminathan, Daiki Kimura, Prithviraj Sen, Keerthiram Murugesan, Rosario Uceda-Sosa, Michiaki Tatsubori, Achille Fokoue, Pavan Kapanipathi, Asim Munawar, et al. 2023. Learning symbolic rules over abstract meaning representations for textual

- reinforcement learning. *arXiv preprint arXiv:2307.02689* (2023).
- [42] Muhao Chen, Yingtao Tian, Mohan Yang, and Carlo Zaniolo. 2016. Multilingual knowledge graph embeddings for cross-lingual knowledge alignment. *arXiv preprint arXiv:1611.03954* (2016).
 - [43] Xinyun Chen, Chen Liang, Adams Wei Yu, Dawn Song, and Denny Zhou. 2020. Compositional generalization via neural-symbolic stack machines. *Advances in Neural Information Processing Systems* 33 (2020), 1690–1701.
 - [44] Xinyun Chen, Chen Liang, Adams Wei Yu, Denny Zhou, Dawn Song, and Quoc V Le. 2019. Neural symbolic reader: Scalable integration of distributed and symbolic representations for reading comprehension. In *International Conference on Learning Representations*.
 - [45] Zhenfang Chen, Rui Sun, Wenjun Liu, Yining Hong, and Chuang Gan. 2023. GENOME: GenerativE Neuro-symbOlic visual reasoning by growing and reusing ModulEs. *arXiv preprint arXiv:2311.04901* (2023).
 - [46] Junyan Cheng and Peter Chin. 2023. On the Transition from Neural Representation to Symbolic Knowledge. *arXiv preprint arXiv:2308.02000* (2023).
 - [47] Rohan Chitnis, Tom Silver, Joshua B Tenenbaum, Tomas Lozano-Perez, and Leslie Pack Kaelbling. 2022. Learning neuro-symbolic relational transition models for bilevel planning. In *2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 4166–4173.
 - [48] Grzegorz Chrupala and Afra Alishahi. 2019. Correlating neural and symbolic representations of language. *arXiv preprint arXiv:1905.06401* (2019).
 - [49] Nuri Cingillioglu. 2022. *End-to-end neuro-symbolic learning of logic-based inference*. Ph.D. Dissertation. Imperial College London.
 - [50] Nuri Cingillioglu and Alessandra Russo. 2021. pix2rule: End-to-end Neuro-symbolic Rule Learning. *arXiv preprint arXiv:2106.07487* (2021).
 - [51] John P Coetzee, Micah A Johnson, Youngzie Lee, Allan D Wu, Marco Iacoboni, and Martin M Monti. 2022. Dissociating language and thought in human reasoning. *Brain Sciences* 13, 1 (2022), 67.
 - [52] William W Cohen, Haitian Sun, R Alex Hofer, and Matthew Siegler. 2020. Scalable neural methods for reasoning with a symbolic knowledge base. *arXiv preprint arXiv:2002.06115* (2020).
 - [53] Matthias Cosler, Christopher Hahn, Ayham Omar, and Frederik Schmitt. 2024. NeuroSynt: A Neuro-symbolic Portfolio Solver for Reactive Synthesis. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 45–67.
 - [54] Miles Cranmer, Alvaro Sanchez Gonzalez, Peter Battaglia, Rui Xu, Kyle Cranmer, David Spergel, and Shirley Ho. 2020. Discovering symbolic models from deep learning with inductive biases. *Advances in neural information processing systems* 33 (2020), 17429–17442.
 - [55] Daniel Cunningham, Mark Law, Jorge Lobo, and Alessandra Russo. 2022. Neuro-symbolic learning of answer set programs from raw data. *arXiv preprint arXiv:2205.12735* (2022).
 - [56] Daniel Cunningham, Mark Law, Jorge Lobo, and Alessandra Russo. 2023. Ffnsl: Feed-forward neural-symbolic learner. *Machine Learning* 112, 2 (2023), 515–569.
 - [57] Matthew L Daggitt, Wen Kokke, Robert Atkey, Natalia Slusarz, Luca Arnaboldi, and Ekaterina Komendantskaya. 2024. Vehicle: Bridging the Embedding Gap in the Verification of Neuro-Symbolic Programs. *arXiv preprint arXiv:2401.06379* (2024).
 - [58] Raphaël Dang-Nhu. 2020. PLANS: Neuro-symbolic program learning from videos. *Advances in Neural Information Processing Systems* 33 (2020), 22445–22455.
 - [59] Alessandro Daniele, Tommaso Campari, Sagar Malhotra, and Luciano Serafini. 2022. Deep symbolic learning: Discovering symbols and rules from perceptions. *arXiv preprint arXiv:2208.11561* (2022).
 - [60] Stéphane d’Ascoli, Pierre-Alexandre Kamienny, Guillaume Lample, and François Charton. 2022. Deep symbolic regression for recurrent sequences. *arXiv preprint arXiv:2201.04600* (2022).
 - [61] Gregory P Davis, Garrett E Katz, Rodolphe J Gentili, and James A Reggia. 2022. NeuroLISP: High-level symbolic programming with attractor neural networks. *Neural Networks* 146 (2022), 200–219.
 - [62] Jacob Devlin, Jonathan Uesato, Rishabh Singh, and Pushmeet Kohli. 2017. Semantic code repair using neuro-symbolic transformation networks. *arXiv preprint arXiv:1710.11054* (2017).
 - [63] Natalia Diaz-Rodriguez, Alberto Lamas, Jules Sanchez, Gianni Franchi, Ivan Donadello, Siham Tabik, David Filliat, Policarpo Cruz, Rosana Montes, and Francisco Herrera. 2022. EXplainable Neural-Symbolic Learning (X-NeSyL) methodology to fuse deep learning representations with expert knowledge graphs: The MonuMAI cultural heritage use case. *Information Fusion* 79 (2022), 58–83.
 - [64] Dominik Dold, Josep Soler Garrido, Victor Caceres Chian, Marcel Hildebrandt, and Thomas Runkler. 2022. Neuro-symbolic computing with spiking neural networks. In *Proceedings of the International Conference on Neuromorphic Systems 2022*. 1–4.
 - [65] Honghua Dong, Jiayuan Mao, Tian Lin, Chong Wang, Lihong Li, and Denny Zhou. 2019. Neural logic machines. *arXiv preprint arXiv:1904.11694* (2019).
 - [66] Paolo Dragone, Stefano Teso, and Andrea Passerini. 2021. Neuro-symbolic constraint programming for structured prediction. *arXiv preprint arXiv:2103.17232* (2021).
 - [67] Rajdeep Dutta, Qincheng Wang, Ankur Singh, Dhruv Kumarjiguda, Li Xiaoli, and Senthilnath Jayavelu. 2023. S-reinforce: A neuro-symbolic policy gradient approach for interpretable reinforcement learning. *arXiv preprint arXiv:2305.07367* (2023).
 - [68] Monireh Ebrahimi, Md Kamruzzaman Sarker, Federico Bianchi, Ning Xie, Aaron Eberhart, Derek Doran, HyeonSik Kim, and Pascal Hitzler. 2021. Neuro-symbolic deductive reasoning for cross-knowledge graph entailment. In *AAAI Spring Symposium: Combining Machine Learning with Knowledge Engineering*.
 - [69] Marco Facchin. 2023. Neural representations unobserved—or: a dilemma for the cognitive neuroscience revolution. *Synthese* 203, 1 (2023), 7.
 - [70] Arnaud Ngumbang Fadjia, Michele Fraccaroli, Alice Bizzarri, Giulia Mazzuchelli, and Evelina Lamma. 2022. Neural-Symbolic Ensemble Learning for early-stage prediction of critical state of Covid-19 patients. *Medical & Biological Engineering & Computing* 60, 12 (2022), 3461–3474.

- [71] Mark Fedyk and Monika Ray. 2023. How to Leverage Machine Learning Interpretability and Explainability to Generate Hypotheses in Cognitive Psychology. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, Vol. 45.
- [72] Reuben Feinman and Brenden M Lake. 2020. Generating new concepts with hybrid neuro-symbolic models. *arXiv preprint arXiv:2003.08978* (2020).
- [73] Reuben Feinman and Brenden M Lake. 2020. Learning task-general representations with generative neuro-symbolic modeling. *arXiv preprint arXiv:2006.14448* (2020).
- [74] Bettina Finzel, Anna Saranti, Alessa Angerschmid, David Tafler, Bastian Pfeifer, and Andreas Holzinger. 2022. Generating explanations for conceptual validation of graph neural networks: An investigation of symbolic predicates learned on relevance-ranked sub-graphs. *KI-Künstliche Intelligenz* 36, 3 (2022), 271–285.
- [75] Amy Sue Fire. 2016. *Learning and Inferring Perceptual Causality from Video*. University of California, Los Angeles.
- [76] João Flach and Luis C Lamb. 2023. A Neural Lambda Calculus: Neurosymbolic AI meets the foundations of computing and functional programming. *arXiv preprint arXiv:2304.09276* (2023).
- [77] Nicholas T Franklin, Kenneth A Norman, Charan Ranganath, Jeffrey M Zacks, and Samuel J Gershman. 2020. Structured Event Memory: A neuro-symbolic model of event cognition. *Psychological Review* 127, 3 (2020), 327.
- [78] P Michael Furlong and Chris Eliasmith. 2023. Modelling neural probabilistic computation using vector symbolic architectures. *Cognitive Neurodynamics* (2023), 1–24.
- [79] Yarin Gal and Zoubin Ghahramani. 2015. Bayesian convolutional neural networks with Bernoulli approximate variational inference. *arXiv preprint arXiv:1506.02158* (2015).
- [80] Artur d’Avila Garcez, Aimore Resende Riquetti Dutra, and Eduardo Alonso. 2018. Towards symbolic reinforcement learning with common sense. *arXiv preprint arXiv:1804.08597* (2018).
- [81] Sankalp Garg, Aniket Bajpai, et al. 2020. Symbolic network: generalized neural policies for relational MDPs. In *International Conference on Machine Learning*. PMLR, 3397–3407.
- [82] Marta Garnelo, Kai Arulkumaran, and Murray Shanahan. 2016. Towards deep symbolic reinforcement learning. *arXiv preprint arXiv:1609.05518* (2016).
- [83] Vedant Gaur and Nikunj Saunshi. 2023. Reasoning in large language models through symbolic math word problems. *arXiv preprint arXiv:2308.01906* (2023).
- [84] Claire Glanois, Zhaohui Jiang, Xuening Feng, Paul Weng, Matthieu Zimmer, Dong Li, Wulong Liu, and Jianye Hao. 2022. Neuro-symbolic hierarchical rule induction. In *International Conference on Machine Learning*. PMLR, 7583–7615.
- [85] Divya Gopinath, Kaiyuan Wang, Mengshi Zhang, Corina S Pasareanu, and Sarfraz Khurshid. 2018. Symbolic execution for deep neural networks. *arXiv preprint arXiv:1807.10439* (2018).
- [86] Zhongyi Han, Benzhen Wei, Xiaoming Xi, Bo Chen, Yilong Yin, and Shuo Li. 2021. Unifying neural learning and symbolic reasoning for spinal medical report generation. *Medical image analysis* 67 (2021), 101872.
- [87] Krishnam Hasija, Shrishti Pradhan, Manasi Patwardhan, Raveendra Kumar Medicherla, Lovekesh Vig, and Ravindra Naik. 2023. Neuro-symbolic Zero-Shot Code Cloning with Cross-Language Intermediate Representation. *arXiv preprint arXiv:2304.13350* (2023).
- [88] Rishi Hazra and Luc De Raedt. 2023. Deep explainable relational reinforcement learning: a neuro-symbolic approach. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*. Springer, 213–229.
- [89] Hao-Yuan He, Wang-Zhou Dai, and Ming Li. 2024. Reduced implication-bias logic loss for neuro-symbolic learning. *Machine Learning* (2024), 1–21.
- [90] Danial Hooshyar. 2024. Temporal learner modelling through integration of neural and symbolic architectures. *Education and Information Technologies* 29, 1 (2024), 1119–1146.
- [91] Danial Hooshyar, Roger Azevedo, and Yeongwook Yang. 2023. Augmenting deep neural networks with symbolic knowledge: Towards trustworthy and interpretable AI for education. *arXiv preprint arXiv:2311.00393* (2023).
- [92] Joy Hsu, Jiayuan Mao, and Jiajun Wu. 2023. Ns3d: Neuro-symbolic grounding of 3d objects and relations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2614–2623.
- [93] Chenxu Hu, Jie Fu, Chenzhuang Du, Simian Luo, Junbo Zhao, and Hang Zhao. 2023. Chatdb: Augmenting llms with databases as their symbolic memory. *arXiv preprint arXiv:2306.03901* (2023).
- [94] Yaojie Hu, Xingjian Shi, Qiang Zhou, and Lee Pike. 2022. Fix bugs with transformer through a neural-symbolic edit grammar. *arXiv preprint arXiv:2204.06643* (2022).
- [95] Ziniu Hu, Yichong Xu, Wenhao Yu, Shuohang Wang, Ziyi Yang, Chenguang Zhu, Kai-Wei Chang, and Yizhou Sun. 2022. Empowering language models with knowledge graph reasoning for question answering. *arXiv preprint arXiv:2211.08380* (2022).
- [96] Jena D Hwang, Chandra Bhagavatula, Ronan Le Bras, Jeff Da, Keisuke Sakaguchi, Antoine Bosselut, and Yejin Choi. 2021. (comet-) atomic 2020: On symbolic and neural commonsense knowledge graphs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 35. 6384–6392.
- [97] León Illanes, Xi Yan, Rodrigo Toro Icarte, and Sheila A McIlraith. [n. d.]. Symbolic plans as high-level instructions for reinforcement learning. In *Proceedings of the international conference on automated planning and scheduling*, Vol. 30. 540–550.
- [98] León Illanes, Xi Yan, Rodrigo Toro Icarte, and Sheila A McIlraith. 2020. Symbolic plans as high-level instructions for reinforcement learning. In *Proceedings of the international conference on automated planning and scheduling*, Vol. 30. 540–550.
- [99] Monika Jain, Kuldeep Singh, and Raghava Mutharaju. 2023. ReOnto: A Neuro-Symbolic Approach for Biomedical Relation Extraction. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*. Springer, 230–247.