

- arXiv:2210.07382 (2022). URL: <https://github.com/cognitiveailab/neurosymbolic>. doi:10.18653/v1/2023.eacl-main.204, submitted 11 February, 2023; v1 submitted 13 October, 2022; originally announced October 2022.
- [43] F. Arabshahi, J. Lee, M. Gawarecki, K. Mazaitis, A. Azaria, T. Mitchell, Conversational neuro-symbolic commonsense reasoning, arXiv preprint arXiv:2006.10022 (2021). URL: <https://github.com/ForoughA/CORGI>. doi:10.1609/aaai.v35i6.16623, submitted 2 February, 2021; v1 submitted 17 June, 2020; originally announced June 2020.
 - [44] D. Chanin, A. Hunter, Neuro-symbolic commonsense social reasoning, arXiv preprint arXiv:2303.08264 (2023). URL: <https://github.com/chanind/amr-logic-converter>, submitted 14 March, 2023; originally announced March 2023.
 - [45] M. Ismayilzada, D. Paul, S. Montariol, M. Geva, A. Bosselut, Crow: Benchmarking commonsense reasoning in real-world tasks, 2023. URL: <https://github.com/mismayil/crow>. doi:10.48550/ARXIV.2310.15239.
 - [46] T. X. Olausson, A. Gu, B. Lipkin, C. E. Zhang, A. Solar-Lezama, J. B. Tenenbaum, R. Levy, Linc: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers, arXiv preprint arXiv:2310.15164 (2023). URL: <https://github.com/benlipkin/linc>. doi:10.18653/v1/2023.emnlp-main.313, submitted 14 February, 2024; v1 submitted 23 October, 2023; originally announced October 2023.
 - [47] J. Prentzas, I. Hatzilygeroudis, Exploring aspects regarding reasoning in neuro-symbolic rules and connectionist expert systems, in: 2021 12th International Conference on Information, Intelligence, Systems & Applications (IISA), 2021, pp. 1–8. URL: <https://github.com/MoserMichael/my-notes/blob/master/talking-with-chatgpt-b.md>. doi:10.1109/IISA52424.2021.9555527.
 - [48] A. Saparov, H. He, Language models are greedy reasoners: A systematic formal analysis of chain-of-thought, 2022. URL: <https://github.com/asaparov/prontoqa>. doi:10.48550/ARXIV.2210.01240.
 - [49] C. Yang, S. Chaudhuri, Safe neurosymbolic learning with differentiable symbolic execution, arXiv preprint arXiv:2203.07671 (2022). URL: <https://github.com/cxyang1997/dse>. doi:10.23919/fmcad.2018.8602991, submitted 15 March, 2022; originally announced March 2022.
 - [50] S. Yang, X. Li, L. Cui, L. Bing, W. Lam, Neuro-symbolic integration brings causal and reliable reasoning proofs, arXiv preprint arXiv:2311.09802 (2023). URL: <https://github.com/damo-nlp-sg/caring>. doi:10.1109/twc.2023.3319981, submitted 16 November, 2023; originally announced November 2023.
 - [51] H. A. Akl, Psychic: A neuro-symbolic framework for knowledge graph question-answering grounding, arXiv preprint arXiv:2310.12638 (2023). URL: <https://github.com/HannaAbiAkl/PSYCHIC>. doi:10.1609/aaai.v35i6.16625, submitted 19 October, 2023; originally announced October 2023.
 - [52] A. Bosselut, R. L. Bras, Y. Choi, Dynamic neuro-symbolic knowledge graph construction for zero-shot commonsense question answering, arXiv preprint arXiv:1911.03876 (2020). URL: <https://github.com/AlirezaDoakhan/QuestionAnswering>. doi:10.1609/aaai.v35i6.16625, submitted 30 October, 2020; v1 submitted 10 November, 2019; originally announced November 2019.
 - [53] S. Cha, H. Oh, Making symbolic execution promising by learning aggressive state-pruning strategy, in: Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering, ESEC/FSE 2020, Association for Computing Machinery, New York, NY, USA, 2020, pp. 147–158. URL: <https://github.com/XMUsun/symbolic-execution-papers>. doi:10.1145/3368089.3409755.
 - [54] S. Chaudhury, P. Sen, M. Ono, D. Kimura, M. Tatsubori, A. Munawar, Neuro-symbolic approaches for text-based policy learning, in: M.-F. Moens, X. Huang, L. Specia, S. W.-t. Yih (Eds.), Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 2021, pp. 3073–3078. URL: <https://github.com/subhajt141/slate-text-based-rl>. doi:10.18653/v1/2021.emnlp-main.245.
 - [55] Z. Chen, R. Sun, W. Liu, Y. Hong, C. Gan, Genome: Generative neuro-symbolic visual reasoning by growing and reusing modules, arXiv preprint arXiv:2311.04901 (2023). URL: <https://github.com/UMass-Foundation-Model/genome>. doi:10.3233/faia230544, submitted 8 November, 2023; originally announced November 2023.
 - [56] R. Chitnis, T. Silver, J. B. Tenenbaum, T. Lozano-Pérez, L. P. Kaelbling, Learning neuro-symbolic relational transition models for bilevel planning, in: 2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), 2022, pp. 4166–4173. URL: <https://github.com/Learning-and-Intelligent-Systems/predicators>. doi:10.1109/IROS47612.2022.9981440.
 - [57] R. Chitnis, T. Silver, J. B. Tenenbaum, T. Lozano-Perez, L. P. Kaelbling, Learning neuro-symbolic relational transition models for bilevel planning, arXiv preprint arXiv:2105.14074 (2022). URL: <https://github.com/Learning-and-Intelligent-Systems/predicators>. doi:10.1109/iros47612.2022.9981440, submitted 30 June, 2022; v1 submitted 28 May, 2021; originally announced May 2021.
 - [58] Y. Feng, X. Yang, X. Zhu, M. Greenspan, Neuro-symbolic natural logic with introspective revision for natural language inference, arXiv preprint arXiv:2203.04857 (2022). URL: <https://github.com/feng-yufei/ns-nli>. doi:10.1162/tac1_a_00458, submitted 5 June, 2022; v1 submitted 9 March, 2022; originally announced March 2022.
 - [59] T. Gupta, A. Kembhavi, Visual programming: Compositional visual reasoning without training, in: 2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 14953–14962. URL: <https://github.com/allenai/visprog>. doi:10.1109/CVPR52729.2023.01436.
 - [60] R. Hazra, B. Chen, A. Rai, N. Kamra, R. Desai, Egotv: Egocentric task verification from natural language task descriptions, in: 2023 IEEE/CVF International Conference on Computer Vision (ICCV), 2023, pp. 15371–15383. URL: <https://github.com/facebookresearch/EgoTV>. doi:10.1109/ICCV51070.2023.01414.
 - [61] R. Hazra, L. D. Raedt, Deep explainable relational reinforcement learning: A neuro-symbolic approach, arXiv preprint arXiv:2304.08349 (2023). URL: <https://github.com/thuwzy/Neural-Symbolic-and-Probabilistic-Logic-Papers>. doi:10.1007/978-3-031-43421-1_13, submitted 14 July, 2023; v1 submitted 17 April, 2023; originally announced April 2023.

- [62] J. He, G. Sivanrupan, P. Tsankov, M. Vechev, Learning to explore paths for symbolic execution, in: Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security, CCS '21, Association for Computing Machinery, New York, NY, USA, 2021, pp. 2526–2540. URL: <https://github.com/eth-sri/learch>. doi:10.1145/3460120.3484813.
- [63] J. Hsu, J. Mao, J. Wu, Ns3d: Neuro-symbolic grounding of 3d objects and relations, in: 2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 2614–2623. URL: <https://github.com/joyhsu0504/NS3D>. doi:10.1109/CVPR52729.2023.00257.
- [64] D. Kimura, M. Ono, S. Chaudhury, R. Kohita, A. Wachi, D. J. Agravante, M. Tatsubori, A. Munawar, A. Gray, Neuro-symbolic reinforcement learning with first-order logic, 2021. URL: <https://github.com/cav-research-lab/safe-reinforcement-learning-using-symbolic-logical-programming-for-autonomous-highway-driving>, <https://ibm.github.io/neuro-symbolic-ai/toolkit/>. doi:10.48550/ARXIV.2110.10963.
- [65] D. Kimura, S. Chaudhury, M. Ono, M. Tatsubori, D. J. Agravante, A. Munawar, A. Wachi, R. Kohita, A. Gray, Loa: Logical optimal actions for text-based interaction games, 2021. URL: <https://github.com/ibm/loa>. doi:10.48550/ARXIV.2110.10973.
- [66] E. van Krieken, T. Thanapalasingam, J. M. Tomczak, F. van Harmelen, A. ten Teije, A-nesi: A scalable approximate method for probabilistic neurosymbolic inference, arXiv preprint arXiv:2212.12393 (2022). URL: <https://github.com/hemile/a-nesi>. doi:10.1007/978-3-540-30215-5_37, submitted 22 September, 2023; v1 submitted 23 December, 2022; originally announced December 2022.
- [67] S. Mghames, L. Castri, M. Hanheide, N. Bellotto, A neuro-symbolic approach for enhanced human motion prediction, in: 2023 International Joint Conference on Neural Networks (IJCNN), 2023, pp. 1–8. URL: <https://github.com/sariahmghames/neurosym-prediction>. doi:10.1109/IJCNN54540.2023.10191970.
- [68] K. Murugesan, M. Atzeni, P. Kapanipathi, P. Shukla, S. Kumaravel, G. Tesauro, K. Talamadupula, M. Sachan, M. Campbell, Text-based rl agents with commonsense knowledge: New challenges, environments and baselines, 2020. URL: <https://github.com/IBM/commonsense-rl>. doi:10.48550/ARXIV.2010.03790.
- [69] T. Nguyen, K. Nguyen, H. Duong, Syminfer: inferring numerical invariants using symbolic states, in: Proceedings of the ACM/IEEE 44th International Conference on Software Engineering: Companion Proceedings, ICSE '22, Association for Computing Machinery, New York, NY, USA, 2022, pp. 197–201. URL: <https://github.com/dynaroars/dig>. doi:10.1145/3510454.3516833.
- [70] K. Sanders, N. Weir, B. V. Durme, Tv-trees: Multimodal entailment trees for neuro-symbolic video reasoning, arXiv preprint arXiv:2402.19467 (2024). URL: <https://github.com/Xuchen-Li/cv-arxiv-daily>. doi:10.18653/v1/2021.emnlp-main.585, submitted 10 March, 2024; v1 submitted 29 February, 2024; originally announced February 2024.
- [71] S. Sanyal, R. K. Manna, K. Roy, Ev-planner: Energy-efficient robot navigation via event-based physics-guided neuromorphic planner, IEEE Robotics and Automation Letters 9 (2024) 2080–2087. URL: <https://github.com/souravsanyal06/ev-planner>. doi:10.1109/LRA.2024.3350982.
- [72] I. Sharifi, M. Yildirim, S. Fallah, Towards safe autonomous driving policies using a neuro-symbolic deep reinforcement learning approach, arXiv preprint arXiv:2307.01316 (2023). URL: <https://github.com/cav-research-lab/safe-reinforcement-learning-using-symbolic-logical-programming-for-autonomous-highway-driving>. doi:10.1109/tpami.2023.3322426, submitted 13 July, 2023; v1 submitted 3 July, 2023; originally announced July 2023.
- [73] T. Silver, A. Athalye, J. B. Tenenbaum, T. Lozano-Perez, L. P. Kaelbling, Learning neuro-symbolic skills for bilevel planning, arXiv preprint arXiv:2206.10680 (2022). URL: <https://github.com/Learning-and-Intelligent-Systems/predicators>. doi:10.1109/iros47612.2022.9981440, submitted 12 October, 2022; v1 submitted 21 June, 2022; originally announced June 2022.
- [74] R. Wang, P. Jansen, M.-A. Côté, P. Ammanabrolu, Scienceworld: Is your agent smarter than a 5th grader?, 2022. URL: <https://github.com/allenai/scienceworld>. doi:10.48550/ARXIV.2203.07540.
- [75] T. Winters, G. Marra, R. Manhaeve, L. De Raedt, Deepstochlog: Neural stochastic logic programming, 2021. URL: <https://github.com/ML-KULeuven/deepstochlog>. doi:10.48550/ARXIV.2106.12574.
- [76] R. Yan, A. Julius, M. Chang, A. Fokoue, T. Ma, R. Uceda-Sosa, Stone: Signal temporal logic neural network for time series classification, in: 2021 International Conference on Data Mining Workshops (ICDMW), 2021, pp. 778–787. URL: <https://github.com/anand-bala/signal-temporal-logic>. doi:10.1109/ICDMW53433.2021.00101.
- [77] R. Yan, T. Ma, A. Fokoue, M. Chang, A. Julius, Neuro-symbolic models for interpretable time series classification using temporal logic description, in: 2022 IEEE International Conference on Data Mining (ICDM), 2022, pp. 618–627. URL: <https://github.com/xiyuanzh/time-series-papers>. doi:10.1109/ICDM54844.2022.00072.
- [78] R. Yan, T. Ma, A. Fokoue, M. Chang, A. Julius, Neuro-symbolic models for interpretable time series classification using temporal logic description, arXiv preprint arXiv:2209.09114 (2022). URL: <https://github.com/xiyuanzh/time-series-papers>. doi:10.1109/icdm54844.2022.00072, submitted 15 September, 2022; originally announced September 2022.
- [79] S. Badreddine, A. d'Avila Garcez, L. Serafini, M. Spranger, Logic tensor networks, Artificial Intelligence 303 (2022) 103649. URL: <https://github.com/logictensornetworks/logictensornetworks>. doi:10.1016/j.artint.2021.103649.
- [80] C. Dickens, C. Gao, C. Pryor, S. Wright, L. Getoor, Convex and bilevel optimization for neuro-symbolic inference and learning, arXiv preprint arXiv:2401.09651 (2024). URL: <https://github.com/convexbilevelnesylearning/psl>. doi:10.1109/iros47612.2022.9981440, submitted 17 January, 2024; originally announced January 2024.
- [81] M. Kodnongbua, L. H. Curtis, A. Schulz, Zero-shot sequential neuro-symbolic reasoning for automatically generating architecture schematic designs, arXiv preprint arXiv:2402.00052 (2024). URL: <https://github.com/Yangyi-Chen/Multimodal-AND-Large-Language-Models>. doi:10.1007/978-3-031-39179-8, submitted 25 January, 2024; originally

- announced February 2024.
- [82] N. Potteiger, X. Koutsoukos, Safe explainable agents for autonomous navigation using evolving behavior trees, in: 2023 IEEE International Conference on Assured Autonomy (ICAA), 2023, pp. 44–52. URL: <https://github.com/tmgthb/Autonomous-Agents>. doi:10.1109/ICAA58325.2023.00014.
 - [83] C. Zhang, B. Jia, S.-C. Zhu, Y. Zhu, Abstract spatial-temporal reasoning via probabilistic abduction and execution, in: 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2021, pp. 9731–9741. URL: <https://github.com/WellyZhang/PrAE>. doi:10.1109/CVPR46437.2021.00961.
 - [84] J. Da, R. L. Bras, X. Lu, Y. Choi, A. Bosselut, Analyzing commonsense emergence in few-shot knowledge models, 2021. URL: <https://github.com/allenai/few-shot-comet>. doi:10.48550/ARXIV.2101.00297.
 - [85] R. Feinman, B. M. Lake, Learning task-general representations with generative neuro-symbolic modeling, arXiv preprint arXiv:2006.14448 (2021). URL: <https://github.com/rfeinman/GNS-Modeling>. doi:10.1145/3539618.3594246, submitted 23 January, 2021; v1 submitted 25 June, 2020; originally announced June 2020.
 - [86] D. Gholipour Ghalandari, C. Hokamp, N. T. Pham, J. Glover, G. Ifrim, A large-scale multi-document summarization dataset from the Wikipedia current events portal, in: D. Jurafsky, J. Chai, N. Schluter, J. Tetreault (Eds.), Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, Association for Computational Linguistics, Online, 2020, pp. 1302–1308. URL: <https://github.com/complementizer/wcep-mds-dataset>. doi:10.18653/v1/2020.acl-main.120. arXiv:<https://aclanthology.org/2020.acl-main.120>.
 - [87] A. Goma, M. Feld, Towards adaptive user-centered neuro-symbolic learning for multimodal interaction with autonomous systems, in: Proceedings of the 25th International Conference on Multimodal Interaction, ICMI '23, Association for Computing Machinery, New York, NY, USA, 2023, pp. 689–694. URL: <https://github.com/tmgthb/Autonomous-Agents>. doi:10.1145/3577190.3616121.
 - [88] M. Yasunaga, J. Leskovec, P. Liang, Linkbert: Pretraining language models with document links, 2022. URL: <https://github.com/michiyasunaga/LinkBERT>. doi:10.48550/ARXIV.2203.15827.
 - [89] E. Zhan, J. J. Sun, A. Kennedy, Y. Yue, S. Chaudhuri, Unsupervised learning of neurosymbolic encoders, arXiv preprint arXiv:2107.13132 (2021). URL: <https://github.com/ezhan94/neurosymbolic-encoders>. doi:10.1016/j.asoc.2019.105851, submitted 20 December, 2022; v1 submitted 27 July, 2021; originally announced July 2021.
 - [90] Y. Zhou, R. Feinman, B. M. Lake, Compositional diversity in visual concept learning, 2023. URL: <https://github.com/yanlizhou/CompositionalDiversity>. doi:10.48550/ARXIV.2305.19374.
 - [91] A. Roberts, H. W. Chung, A. Levskaya, G. Mishra, J. Bradbury, D. Andor, S. Narang, B. Lester, C. Gaffney, A. Mohiuddin, C. Hawthorne, A. Lewkowycz, A. Salcianu, M. van Zee, J. Austin, S. Goodman, L. B. Soares, H. Hu, S. Tsvyashchenko, A. Chowdhery, J. Bastings, J. Bulian, X. Garcia, J. Ni, A. Chen, K. Kenealy, J. H. Clark, S. Lee, D. Garrette, J. Lee-Thorp, C. Raffel, N. Shazeer, M. Ritter, M. Bosma, A. Passos, J. Maitin-Shepard, N. Fiedel, M. Omernick, B. Saeta, R. Sepassi, A. Spiridonov, J. Newlan, A. Gesmundo, Scaling up models and data with t5x and seqio, 2022. URL: <https://github.com/google-research/t5x>. doi:10.48550/ARXIV.2203.17189.
 - [92] F. Fabiano, V. Pallagani, M. B. Ganapini, L. Horeish, A. Loreggia, K. Murugesan, F. Rossi, B. Srivastava, Plan-SOFAI: A neuro-symbolic planning architecture, in: Neuro-Symbolic Learning and Reasoning in the era of Large Language Models, 2023. URL: <https://github.com/FrancescoFabiano/SOFAI-x-Planning>. arXiv:<https://openreview.net/forum?id=ORAhay0H4x>.
 - [93] D. J. Agravante, D. Kimura, M. Tatsubori, Learning neuro-symbolic world models with logical neural networks, in: PRL Workshop Series – Bridging the Gap Between AI Planning and Reinforcement Learning, 2023. URL: <https://github.com/IBM/LNN>. arXiv:<https://openreview.net/forum?id=VD0ksRT1jb>.
 - [94] N. Haut, W. Banzhaf, B. Punch, Active learning improves performance on symbolic regression tasks in stackgp, in: Proceedings of the Genetic and Evolutionary Computation Conference Companion, GECCO '22, Association for Computing Machinery, New York, NY, USA, 2022, pp. 550–553. URL: <https://github.com/hoolagans/stackgp>. doi:10.1145/3520304.3528941.
 - [95] N. Haut, B. Punch, W. Banzhaf, Active learning informs symbolic regression model development in genetic programming, in: Proceedings of the Companion Conference on Genetic and Evolutionary Computation, GECCO '23 Companion, Association for Computing Machinery, New York, NY, USA, 2023, pp. 587–590. URL: <https://github.com/hoolagans/stackgp>. doi:10.1145/3583133.3590577.
 - [96] J. R. Kirk, R. E. Wray, J. E. Laird, Exploiting language models as a source of knowledge for cognitive agents, Proceedings of the AAAI Symposium Series 2 (2024) 286–294. URL: <https://github.com/zjunlp/Prompt4ReasoningPapers>. doi:10.1609/aaais.v2i1.27690.
 - [97] K. Knowles, M. Witbrock, G. Dobbie, V. Yogarajan, A proposal for a language model based cognitive architecture, Proceedings of the AAAI Symposium Series 2 (2024) 295–301. URL: <https://github.com/Aryia-Behroziuan/References>. doi:10.1609/aaais.v2i1.27691.
 - [98] S. Kulal, J. Mao, A. Aiken, J. Wu, Hierarchical motion understanding via motion programs, in: 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2021, pp. 6564–6572. URL: https://github.com/Sumith1896/motion2prog_release. doi:10.1109/CVPR46437.2021.00650.
 - [99] A. Liu, S. Swayamdipta, N. A. Smith, Y. Choi, Wanli: Worker and ai collaboration for natural language inference dataset creation, 2022. URL: <https://github.com/alisaawuffles/wanli>. doi:10.48550/ARXIV.2201.05955.
 - [100] A. Ross, T. Wu, H. Peng, M. Peters, M. Gardner, Tailor: Generating and perturbing text with semantic controls, in: Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Association for Computational Linguistics, 2022. URL: <https://github.com/allenai/tailor>. doi:10.18653/v1/2022.acl-long.228.