

- [101] A. Smirnova, J. Yang, D. Yang, P. Cudre-Mauroux, Nesy: A neuro-symbolic system for label noise reduction, *IEEE Transactions on Knowledge and Data Engineering* 35 (2023) 8300–8311. URL: <https://github.com/eXascaleInfolab/Nesy>. doi:10.1109/TKDE.2022.3199570.
- [102] e. a. Srivastava, Beyond the imitation game: Quantifying and extrapolating the capabilities of language models (2022). URL: <https://github.com/google/BIG-bench>. doi:10.48550/ARXIV.2206.04615.
- [103] S. Sukhbaatar, D. JU, S. Poff, S. Roller, A. Szlam, J. E. Weston, A. Fan, Not all memories are created equal: Learning to expire, 2021. URL: <https://github.com/lucidrains/learning-to-expire-pytorch>. doi:10.1037/e527342012-077. arXiv:https://openreview.net/forum?id=ZVBtN6B_6i7.
- [104] Y. Wu, M. Gardner, P. Stenetorp, P. Dasigi, Generating data to mitigate spurious correlations in natural language inference datasets, in: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics, 2022. URL: <https://github.com/jimmycode/gen-debiased-nli>. doi:10.18653/v1/2022.acl-long.190.
- [105] J. D. Hwang, C. Bhagavatula, R. Le Bras, J. Da, K. Sakaguchi, A. Bosselut, Y. Choi, (comet-) atomic 2020: On symbolic and neural commonsense knowledge graphs, *Proceedings of the AAAI Conference on Artificial Intelligence* 35 (2021) 6384–6392. URL: <https://github.com/allenai/comet-atomic-2020>. doi:10.1609/aaai.v35i7.16792.
- [106] M. Ismayilzada, A. Bosselut, kogito: A commonsense knowledge inference toolkit, 2022. URL: <https://github.com/epfl-nlp/kogito>. doi:10.48550/ARXIV.2211.08451.
- [107] N. Mostafazadeh, A. Kalyanpur, L. Moon, D. Buchanan, L. Berkowitz, O. Biran, J. Chu-Carroll, Glucose: Generalized and contextualized story explanations, 2020. URL: <https://github.com/ElementalCognition/glucose>. doi:10.48550/ARXIV.2009.07758.
- [108] N. Papoulias, Neuroql: A neuro-symbolic language and dataset for inter-subjective reasoning, *arXiv preprint arXiv:2303.07146* (2023). URL: <https://github.com/orgdlabs/neuroQL>. doi:10.1109/cict56698.2022.9997814, submitted 13 March, 2023; originally announced March 2023.
- [109] A. Perevalov, D. Diefenbach, R. Usbeck, A. Both, Qald-9-plus: A multilingual dataset for question answering over dbpedia and wikidata translated by native speakers, 2022. URL: https://github.com/perevalov/qald_9_plus. doi:10.48550/ARXIV.2202.00120.
- [110] D. N. Ribeiro, K. Forbus, Combining analogy with language models for knowledge extraction, in: *3rd Conference on Automated Knowledge Base Construction*, 2021. URL: <https://github.com/dnr2/analogical-ke>. doi:10.59350/mk9dg-4r198.
- [111] Z. Chen, G. Weiss, E. Mitchell, A. Celikyilmaz, A. Bosselut, Reckoning: Reasoning through dynamic knowledge encoding, 2023. URL: <https://github.com/eric11eca/reckoning-metaskg>. doi:10.48550/ARXIV.2305.06349.
- [112] A. Agafonov, A. Ponomarev, Localization of ontology concepts in deep convolutional neural networks, in: *2022 IEEE International Multi-Conference on Engineering, Computer and Information Sciences (SIBIRCON)*, 2022, pp. 160–165. URL: <https://github.com/aaagafonov/Concept-Localization-Experiments>. doi:10.1109/SIBIRCON56155.2022.10016932.
- [113] I. Harmon, S. Marconi, B. Weinstein, Y. Bai, D. Z. Wang, E. White, S. Bohlman, Improving rare tree species classification using domain knowledge, *IEEE Geoscience and Remote Sensing Letters* 20 (2023) 1–5. URL: <https://github.com/satellite-image-deep-learning/techniques>. doi:10.1109/LGRS.2023.3278170.
- [114] T. Racharak, On approximation of concept similarity measure in description logic elh with pre-trained word embedding, *IEEE Access* 9 (2021) 61429–61443. URL: <https://github.com/monologg/nlp-arxiv-daily>. doi:10.1109/ACCESS.2021.3073730.
- [115] K. Raj, A neuro-symbolic approach to enhance interpretability of graph neural network through the integration of external knowledge, in: *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management, CIKM '23*, Association for Computing Machinery, New York, NY, USA, 2023, pp. 5177–5180. URL: <https://github.com/thuwzy/Neural-Symbolic-and-Probabilistic-Logic-Papers>. doi:10.1145/3583780.3616008.
- [116] L. F. R. Ribeiro, M. Liu, I. Gurevych, M. Dreyer, M. Bansal, Factgraph: Evaluating factuality in summarization with semantic graph representations, 2022. URL: <https://github.com/amazon-research/fact-graph>. doi:10.48550/ARXIV.2204.06508.
- [117] J. Chen, D. Yang, Structure-aware abstractive summarization via discourse and action graphs, 2021. URL: <https://github.com/GT-SALT/Structure-Aware-BART>. doi:10.48550/ARXIV.2104.08400.
- [118] J. Hessel, A. Marasović, J. D. Hwang, L. Lee, J. Da, R. Zellers, R. Mankoff, Y. Choi, Do androids laugh at electric sheep? humor "understanding" benchmarks from the new yorker caption contest (2022). URL: https://github.com/jmhessel/caption_contest_corpus. doi:10.48550/ARXIV.2209.06293.
- [119] A. Kalyanpur, T. Breloff, D. Ferrucci, Braid: Weaving symbolic and neural knowledge into coherent logical explanations, 2020. URL: <https://github.com/chanind/tensor-theorem-prover>. doi:10.48550/ARXIV.2011.13354.
- [120] D. Wan, M. Bansal, Factpegasus: Factuality-aware pre-training and fine-tuning for abstractive summarization, 2022. URL: <https://github.com/meetdavidwan/factpegasus>. doi:10.48550/ARXIV.2205.07830.
- [121] F. Künnecke, A. Filighera, C. Leong, T. Steuer, Enhancing multi-domain automatic short answer grading through an explainable neuro-symbolic pipeline, *arXiv preprint arXiv:2403.01811* (2024). URL: <https://github.com/iwangjian/Paper-Reading-ConvAI>. doi:10.18653/v1/2023.bea-1.29, submitted 19 March, 2024; v1 submitted 4 March, 2024; originally announced March 2024.
- [122] K. Ahmed, K.-W. Chang, G. V. d. Broeck, A pseudo-semantic loss for autoregressive models with logical constraints, 2023. URL: <https://github.com/UCLA-StarAI/PseudoSL>. doi:10.48550/ARXIV.2312.03905.
- [123] A. Ahmetoglu, M. Y. Seker, J. Piater, E. Oztop, E. Ugur, Deepsym: Deep symbol generation and rule learning for planning from unsupervised robot interaction, *J. Artif. Int. Res.* 75 (2022). URL: <https://github.com/alper111/DeepSym>. doi:10.1613/jair.1.13754.

- [124] K. G. Baugh, N. Cingillioglu, A. Russo, Neuro-symbolic rule learning in real-world classification tasks, arXiv preprint arXiv:2303.16674 (2023). URL: <https://github.com/kittykg/neural-dnf-cub>. doi:10.18653/v1/2023.acl-short.57, submitted 29 March, 2023; originally announced March 2023.
- [125] R. Calinescu, C. Imrie, R. Mangal, G. N. Rodrigues, C. Păsăreanu, M. A. Santana, G. Vázquez, Controller synthesis for autonomous systems with deep-learning perception components, IEEE Transactions on Software Engineering (2024) 1–22. URL: <https://github.com/ccimrie/DeepDECS>. doi:10.1109/TSE.2024.3385378.
- [126] N. Cingillioglu, A. Russo, pix2rule: End-to-end neuro-symbolic rule learning, arXiv preprint arXiv:2106.07487 (2021). URL: <https://github.com/kittykg/neural-dnf-cub>. doi:10.1007/978-3-030-86340-1_4, submitted 28 February, 2022; v1 submitted 14 June, 2021; originally announced June 2021.
- [127] D. Cunningham, M. Law, J. Lobo, A. Russo, Neuro-symbolic learning of answer set programs from raw data, arXiv preprint arXiv:2205.12735 (2022). URL: <https://github.com/dancunnington/nsil>. doi:10.24963/ijcai.2023/399, submitted 2 February, 2024; v1 submitted 25 May, 2022; originally announced May 2022.
- [128] D. Dold, J. S. Garrido, An energy-based model for neuro-symbolic reasoning on knowledge graphs, in: 2021 20th IEEE International Conference on Machine Learning and Applications (ICMLA), 2021, pp. 916–921. URL: <https://github.com/dodo47/cyberml>. doi:10.1109/ICMLA52953.2021.00151.
- [129] S. Gao, J. D. Hwang, S. Kanno, H. Wakaki, Y. Mitsufuji, A. Bosselut, Comfact: A benchmark for linking contextual commonsense knowledge, 2022. URL: <https://github.com/silin159/comfact>. doi:10.48550/ARXIV.2210.12678.
- [130] S. Gao, B. Borges, S. Oh, D. Bayazit, S. Kanno, H. Wakaki, Y. Mitsufuji, A. Bosselut, Peacok: Persona commonsense knowledge for consistent and engaging narratives, 2023. URL: <https://github.com/silin159/peacok>. doi:10.48550/ARXIV.2305.02364.
- [131] I. Harmon, S. Marconi, B. Weinstein, S. Graves, D. Z. Wang, A. Zare, S. Bohlman, A. Singh, E. White, Injecting domain knowledge into deep neural networks for tree crown delineation, IEEE Transactions on Geoscience and Remote Sensing 60 (2022) 1–19. URL: <https://github.com/satellite-image-deep-learning/techniques>. doi:10.1109/TGRS.2022.3216622.
- [132] T. He, L. Gao, J. Song, Y.-F. Li, Toward a unified transformer-based framework for scene graph generation and human-object interaction detection, IEEE Transactions on Image Processing 32 (2023) 6274–6288. URL: <https://github.com/DirtyHarryLYL/Transformer-in-Vision>. doi:10.1109/TIP.2023.3330304.
- [133] J. Hong, T. P. Pavlic, An insect-inspired randomly, weighted neural network with random fourier features for neuro-symbolic relational learning, arXiv preprint arXiv:2109.06663 (2021). URL: <https://github.com/jyhong0304/SII>. doi:10.1007/11816157_123, submitted 11 September, 2021; originally announced September 2021.
- [134] P. Howard, J. Wang, V. Lal, G. Singer, Y. Choi, S. Swayamdipta, Neurocomparatives: Neuro-symbolic distillation of comparative knowledge, arXiv preprint arXiv:2305.04978 (2023). URL: <https://github.com/KSESEU/LLMPapers/blob/main/README.md>. doi:10.5220/0011718500003393, submitted 5 April, 2024; v1 submitted 8 May, 2023; originally announced May 2023.
- [135] J. Huang, Z. Li, M. Naik, S.-N. Lim, Laser: A neuro-symbolic framework for learning spatial-temporal scene graphs with weak supervision, arXiv preprint arXiv:2304.07647 (2023). URL: <https://github.com/isLinXu/paper-list>. doi:10.23919/eusipco58844.2023.10289877, submitted 22 November, 2023; v1 submitted 15 April, 2023; originally announced April 2023.
- [136] A. Khatiwada, S. Shirai, K. Srinivas, O. Hassanzadeh, Knowledge graph embeddings for causal relation prediction, in: Workshop on Deep Learning for Knowledge Graphs (DL4KG@ISWC2022), CEUR Workshop Proceedings, CEUR-WS.org, Yorktown Heights, NY, USA; Boston, MA, USA; Troy, NY, USA, 2022. URL: <https://github.com/zjukg/cause>. doi:10.1088/1742-6596/1871/1/012050. arXiv:https://ceur-ws.org/Vol-3342/paper-8.pdf, work done while at IBM Research.
- [137] H. Kim, J. Bak, K. Cho, H. Koo, A transformer-based function symbol name inference model from an assembly language for binary reversing, in: Proceedings of the 2023 ACM Asia Conference on Computer and Communications Security, ASIA CCS '23, Association for Computing Machinery, New York, NY, USA, 2023, pp. 951–965. URL: <https://github.com/agwaBom/AsmDepictor>. doi:10.1145/3579856.3582823.
- [138] T. Klinger, L. Liu, S. Dan, M. Crouse, P. Ram, A. Gray, Compositional program generation for few-shot systematic generalization, 2023. URL: <https://github.com/IBM/cpg>. doi:10.48550/ARXIV.2309.16467.
- [139] Y. Li, J. Mao, X. Zhang, W. T. Freeman, J. B. Tenenbaum, J. Wu, Perspective plane program induction from a single image, in: 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020, pp. 4433–4442. URL: <https://github.com/42x00/p3i>. doi:10.1109/CVPR42600.2020.00449.
- [140] H. Liang, W. Lei, P. Y. Chan, Z. Yang, M. Sun, T.-S. Chua, Pirhdy: Learning pitch-, rhythm-, and dynamics-aware embeddings for symbolic music, in: Proceedings of the 28th ACM International Conference on Multimedia, MM '20, Association for Computing Machinery, New York, NY, USA, 2020, pp. 574–582. URL: <https://github.com/mengshor/PiRhDy>. doi:10.1145/3394171.3414032.
- [141] L. Luo, G. Zhang, H. Xu, Y. Yang, C. Fang, Q. Li, Insight: End-to-end neuro-symbolic visual reinforcement learning with language explanations, arXiv preprint arXiv:2403.12451 (2024). URL: <https://github.com/liruilo/nsrl-vision-pub>. doi:10.18653/v1/2022.findings-emnlp.345, submitted 19 March, 2024; originally announced March 2024.
- [142] E. Marconato, G. Bontempo, E. Ficarra, S. Calderara, A. Passerini, S. Teso, Neuro-symbolic continual learning: Knowledge, reasoning shortcuts and concept rehearsal, arXiv preprint arXiv:2302.01242 (2023). URL: <https://github.com/ema-marconato/nesy-cl>. doi:10.1007/978-3-031-39179-8, submitted 19 December, 2023; v1 submitted 2 February, 2023; originally announced February 2023.
- [143] E. Marconato, S. Teso, A. Vergari, A. Passerini, Not all neuro-symbolic concepts are created equal: Analysis and mitigation of

- reasoning shortcuts, arXiv preprint arXiv:2305.19951 (2023). URL: <https://github.com/ema-marconato/reasoning-shortcuts>. doi:10.1017/cbo9780511574917.007, submitted 18 December, 2023; v1 submitted 31 May, 2023; originally announced May 2023.
- [144] E. Marconato, S. Bortolotti, E. van Krieken, A. Vergari, A. Passerini, S. Teso, Bears make neuro-symbolic models aware of their reasoning shortcuts, arXiv preprint arXiv:2402.12240 (2024). URL: <https://github.com/samuelebortolotti/bears>. doi:10.1007/978-3-031-39179-8, submitted 19 February, 2024; originally announced February 2024.
 - [145] S. Martone, F. Manigrasso, F. Lamberti, L. Morra, Prototypical logic tensor networks (proto-ltn) for zero shot learning, in: 2022 26th International Conference on Pattern Recognition (ICPR), 2022, pp. 4427–4433. URL: <https://github.com/francescomanigrasso/proto-ltn>. doi:10.1109/ICPR56361.2022.9956239.
 - [146] K. Mukherji, D. Parkar, L. Pokala, D. Aditya, P. Shakarian, C. Dorman, Scalable semantic non-markovian simulation proxy for reinforcement learning, in: 2024 IEEE 18th International Conference on Semantic Computing (ICSC), 2024, pp. 183–190. URL: <https://github.com/yingchengyang/Reinforcement-Learning-Papers>. doi:10.1109/ICSC59802.2024.00035.
 - [147] N. K. H. Singh, V. Bindal, A. Tuli, V. Agrawal, R. Jain, P. Singla, R. Paul, Learning neuro-symbolic programs for language guided robot manipulation, in: 2023 IEEE International Conference on Robotics and Automation (ICRA), 2023, pp. 7973–7980. URL: <https://github.com/dair-iitd/nsrmp>. doi:10.1109/ICRA48891.2023.10160545.
 - [148] C. Núñez-Molina, P. Mesejo, J. Fernández-Olivares, Nesig: A neuro-symbolic method for learning to generate planning problems, arXiv preprint arXiv:2301.10280 (2023). URL: <https://github.com/IBM/neuro-symbolic-ai>. doi:10.1609/aaai.v38i21.30409, submitted 24 January, 2023; originally announced January 2023.
 - [149] S. Oruganti, S. Nirenburg, J. English, M. McShane, Automating knowledge acquisition for content-centric cognitive agents using llms, Proceedings of the AAAI Symposium Series 2 (2024) 379–385. URL: <https://github.com/AGI-Edgerunners/LLM-Agents-Papers>. doi:10.1609/aaais.v2i1.27703.
 - [150] A. Potnis, D. Lunga, A. Sorokine, P. Dias, L. Yang, J. Arndt, J. Bowman, J. Wohlgemuth, Towards geospatial knowledge graph infused neuro-symbolic ai for remote sensing scene understanding, in: IGARSS 2023 - 2023 IEEE International Geoscience and Remote Sensing Symposium, 2023, pp. 1400–1403. URL: <https://github.com/isLinXu/paper-list>. doi:10.1109/IGARSS52108.2023.10281958.
 - [151] D. Purohit, Y. Chudasama, A. Rivas, M.-E. Vidal, Sparkle: Symbolic capturing of knowledge for knowledge graph enrichment with learning, in: Proceedings of the 12th Knowledge Capture Conference 2023, K-CAP '23, Association for Computing Machinery, New York, NY, USA, 2023, pp. 44–52. URL: <https://github.com/SDM-TIB/SPARKLE>. doi:10.1145/3587259.3627547.
 - [152] C. Raymond, Q. Chen, B. Xue, M. Zhang, Learning symbolic model-agnostic loss functions via meta-learning, IEEE Transactions on Pattern Analysis and Machine Intelligence 45 (2023) 13699–13714. URL: <https://github.com/decadz/evolved-model-agnostic-loss>. doi:10.1109/TPAMI.2023.3294394.
 - [153] N. Ruaro, K. Zeng, L. Dresel, M. Polino, T. Bao, A. Continella, S. Zanero, C. Kruegel, G. Vigna, Syml: Guiding symbolic execution toward vulnerable states through pattern learning, in: Proceedings of the 24th International Symposium on Research in Attacks, Intrusions and Defenses, RAID '21, Association for Computing Machinery, New York, NY, USA, 2021, pp. 456–468. URL: <https://github.com/XMUusny/symbolic-execution-papers>. doi:10.1145/3471621.3471865.
 - [154] S. S. Saha, S. S. Sandha, M. Aggarwal, B. Wang, L. Han, J. D. G. Briseno, M. Srivastava, Tinyns: Platform-aware neurosymbolic auto tiny machine learning, ACM Trans. Embed. Comput. Syst. 23 (2024). URL: <https://github.com/nesl/neurosymbolic-tinym>. doi:10.1145/3603171.
 - [155] S. Shirai, D. Bhattacharjya, O. Hassanzadeh, Event prediction using case-based reasoning over knowledge graphs, in: Proceedings of the ACM Web Conference 2023, WWW '23, ACM, 2023. URL: <https://github.com/solashirai/www-evcbr>. doi:10.1145/3543507.3583201.
 - [156] L. Tao, Y.-X. Huang, W.-Z. Dai, Y. Jiang, Deciphering raw data in neuro-symbolic learning with provable guarantees, arXiv preprint arXiv:2308.10487 (2023). URL: <https://github.com/abductivelearning/abl-tl>. doi:10.1609/aaai.v38i14.29455, submitted 23 January, 2024; v1 submitted 21 August, 2023; originally announced August 2023.
 - [157] A. Valenti, D. Bacciu, Leveraging relational information for learning weakly disentangled representations, in: 2022 International Joint Conference on Neural Networks (IJCNN), 2022, pp. 1–8. URL: <https://github.com/andrea-v/weak-disentanglement>. doi:10.1109/IJCNN55064.2022.9892093.
 - [158] Y. Wang, Z. Tu, Y. Xiang, S. Zhou, X. Chen, B. Li, T. Zhang, Rapid image labeling via neuro-symbolic learning, in: Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD '23, Association for Computing Machinery, New York, NY, USA, 2023, pp. 2467–2477. URL: <https://github.com/Neural-Symbolic-Image-Labeling/Rapid>. doi:10.1145/3580305.3599485.
 - [159] L. Werner, N. Layaída, P. Genevès, S. Chlyah, Knowledge enhanced graph neural networks, in: 2023 IEEE 10th International Conference on Data Science and Advanced Analytics (DSAA), 2023, pp. 1–10. URL: <https://github.com/RUCDM/KB4Rec>. doi:10.1109/DSAA60987.2023.10302495.
 - [160] T. Wu, M. Tjandrasuwita, Z. Wu, X. Yang, K. Liu, R. Sosič, J. Leskovec, Zeroc: A neuro-symbolic model for zero-shot concept recognition and acquisition at inference time, arXiv preprint arXiv:2206.15049 (2022). URL: <https://github.com/snap-stanford/zeroc>. doi:10.1109/tmm.2023.3318300, submitted 11 October, 2022; v1 submitted 30 June, 2022; originally announced June 2022.
 - [161] K. Ahmed, S. Teso, K.-W. Chang, G. V. den Broeck, A. Vergari, Semantic probabilistic layers for neuro-symbolic learning, arXiv preprint arXiv:2206.00426 (2022). URL: <https://github.com/thuwzy/Neural-Symbolic-and-Probabilistic-Logic-Papers>. doi:10.3233/faia230153, submitted 1 June, 2022; originally announced June 2022.